

sheet6

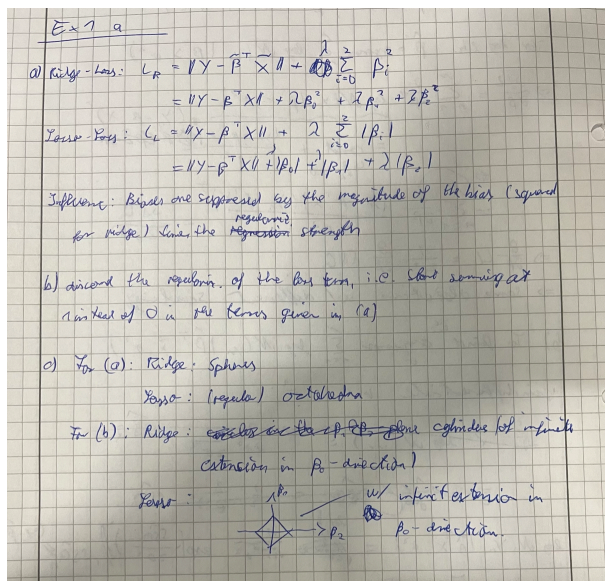
December 5, 2022

1 Sheet 6

Leonard Benkendorff, Dorothea Schwärzel, Sebastian Preuß

```
[1]: import numpy as np
from matplotlib import pyplot as plt
import scipy.sparse
from sklearn.linear_model import Ridge, LinearRegression
```

1.1 1 Regularization and Bias



1.2 2 Estimating Parameter Relevance

```
[2]: # load the data
with open('data/vostok.txt', 'r') as f:
    lines = f.readlines()

# remove header and split lines
lines = [l.split() for l in lines[2:]]
```

```

# filter out lines with missing data
lines = [l for l in lines if len(l) == 4]

# convert to float
lines = np.array(lines).astype(np.float32)
print(f'{lines.shape=}')

features = np.concatenate([lines[:, :1], lines[:, 2:]], axis=1).T
feature_names = 'age', 'CO ', 'dust'
labels = lines[:, 1]
label_name = 'T'

print(f'{features.shape=}, {labels.shape=}')

```

```

lines.shape=(3729, 4)
features.shape=(3, 3729), labels.shape=(3729,)

```

```

[3]: from sklearn.linear_model import LinearRegression

# TODO: fit the linear regressor and compute the sum of square deviations
lr = LinearRegression()
lr.fit(features.T, labels)

print('Sum of squared deviations:', np.sum((lr.predict(features.T)-labels)**2))


plt.plot(labels, label='True data')
plt.plot(lr.predict(features.T), label='Predicted data')
plt.legend(loc='best')
plt.title(r'True and predicted values for $\Delta T$')
plt.show()

```

```

Sum of squared deviations: 6362.9375

```



```
[4]: # TODO: for each feature, randomly permute it amongst the samples,
#       refit the regressor and compute sum of squared deviations
for i in range(3):
    new_order = np.arange(features.shape[-1])
    np.random.shuffle(new_order)
    shuffled_features = np.zeros(features.shape)
    for j in range(3):
        if i == j:
            shuffled_features[i,:] = features[i,new_order]
        else:
            shuffled_features[j,:] = features[j,:]


    lr = LinearRegression()
    lr.fit(shuffled_features.T, labels)
    print(f'Sum of squared deviations after shuffling {feature_names[i]}:', np.
    ↪sum((lr.predict(features.T)-labels)**2))
    plt.plot(lr.predict(shuffled_features.T), label=f'Prediction after_
    ↪shuffling {feature_names[i]}')
plt.plot(labels, label='True data', color='#8c564baa') # add some transparency
plt.legend(loc='best')
```

```
plt.title(r'True and predicted values for $\Delta T$ after shuffling')
plt.show()
```

Sum of squared deviations after shuffling age: 6796.764275972411

Sum of squared deviations after shuffling CO₂ : 18962.56096840998

Sum of squared deviations after shuffling dust: 6565.202563606381



Which feature is most important and which is least relevant? The difference to the original sum of squared deviations is as follows:

$$\Delta_{\text{age}} = |6362.9375 - 6824.1124| = 461.2 \quad (1)$$

$$\Delta_{\text{CO}_2} = |6362.9375 - 18944.9081| = 12582.0 \quad (2)$$

$$\Delta_{\text{dust}} = |6362.9375 - 6537.4003| = 174.5 \quad (3)$$

(This may differ from the output above if the code is re-run, because `np.random.shuffle` returns a new, random, result every time).

We conclude: The feature CO₂ is the most important one, the feature dust is the least important one. This conclusion can (at least for CO₂) also be drawn from the visualisation we included.

1.3 3 σ^2 -estimation and Heteroscedastic Noise

ML sheet 06

a) $SSQ = (Y - \beta^T X)^T (Y - \beta^T X)^T$
 $\frac{\partial SSQ}{\partial \beta} = -2 Y X^T + 2 X^T X \beta = 0$
 $\Rightarrow Y X^T = \beta^T X X^T, \hat{\beta} = (X X^T)^{-1} Y X^T$
 for log lik. $\hat{\beta} = \arg \max_{\beta} \sum_{n=1}^N \log N(y_n | \beta^T x_n, \sigma^2)$
 $\frac{\partial LL}{\partial \beta} = \sum_{n=1}^N \frac{\partial}{\partial \beta} \log N(y_n | \beta^T x_n, \sigma^2)$
 $= \sum_{n=1}^N \frac{1}{\sigma^2} \frac{\partial}{\partial \beta} \left(-\frac{(y_n - \beta^T x_n)^2}{2\sigma^2} \right)$
 $= \sum_{n=1}^N \frac{1}{\sigma^2} \frac{\partial}{\partial \beta} \left(-\frac{(y_n - \beta^T x_n)^2}{2\sigma^2} \right) = \frac{SSQ}{2\sigma^2}$
 $\Rightarrow \frac{\partial LL}{\partial \beta} = \frac{1}{2\sigma^2} \frac{\partial}{\partial \beta} SSQ$
 \Rightarrow same direction w/ different scaling

b) to find $\hat{\sigma}^2$ $\sum_{n=1}^N \log N(y_n | \hat{\beta}^T x_n, \sigma^2)$
 again differentiate w.r.t. σ^2 :
 $\frac{\partial}{\partial \sigma^2} \sum_{n=1}^N \log N(y_n | \hat{\beta}^T x_n, \sigma^2) = \sum_{n=1}^N \frac{\partial}{\partial \sigma^2} \left(-\frac{(y_n - \hat{\beta}^T x_n)^2}{2\sigma^2} - \frac{1}{2} \log \sigma^2 \right)$
 $= \sum_{n=1}^N \left(\frac{(y_n - \hat{\beta}^T x_n)^2}{2\sigma^4} - \frac{1}{2\sigma^2} \right)$
 $= \frac{1}{2\sigma^4} \sum_{n=1}^N (y_n - \hat{\beta}^T x_n)^2 - \frac{N}{2\sigma^2}$
 $= \frac{1}{2\sigma^4} SSQ - \frac{N}{2\sigma^2}$

Comparison to SSQ: In the lecture we derived the covariance of β from the distribution of the variance of the original data. We did not find an expression for the likeliest variance σ^2 .

1.4 4 Visualize Regularization Contours

```
[5]: # load the data
data = np.load('data/linreg.npz')
x = data['X']
y = data['Y']
print(f'{x.shape} {y.shape}')
```

(2, 100) (1, 100)

```
[6]: # TODO: create a grid of points in the parameter space
beta1 = np.linspace(-1, 3, 150)
beta2 = np.linspace(-1, 3, 150)
beta1, beta2 = np.meshgrid(beta1, beta2)
```

(a)

```
[7]: # TODO: make contour plots for ridge and lasso regularization terms
def regulariz_ridge(beta1, beta2): # essentially the 2-norm
    return beta1*beta1 + beta2*beta2

def regulariz_lasso(beta1, beta2): # essentially the 1-norm
```

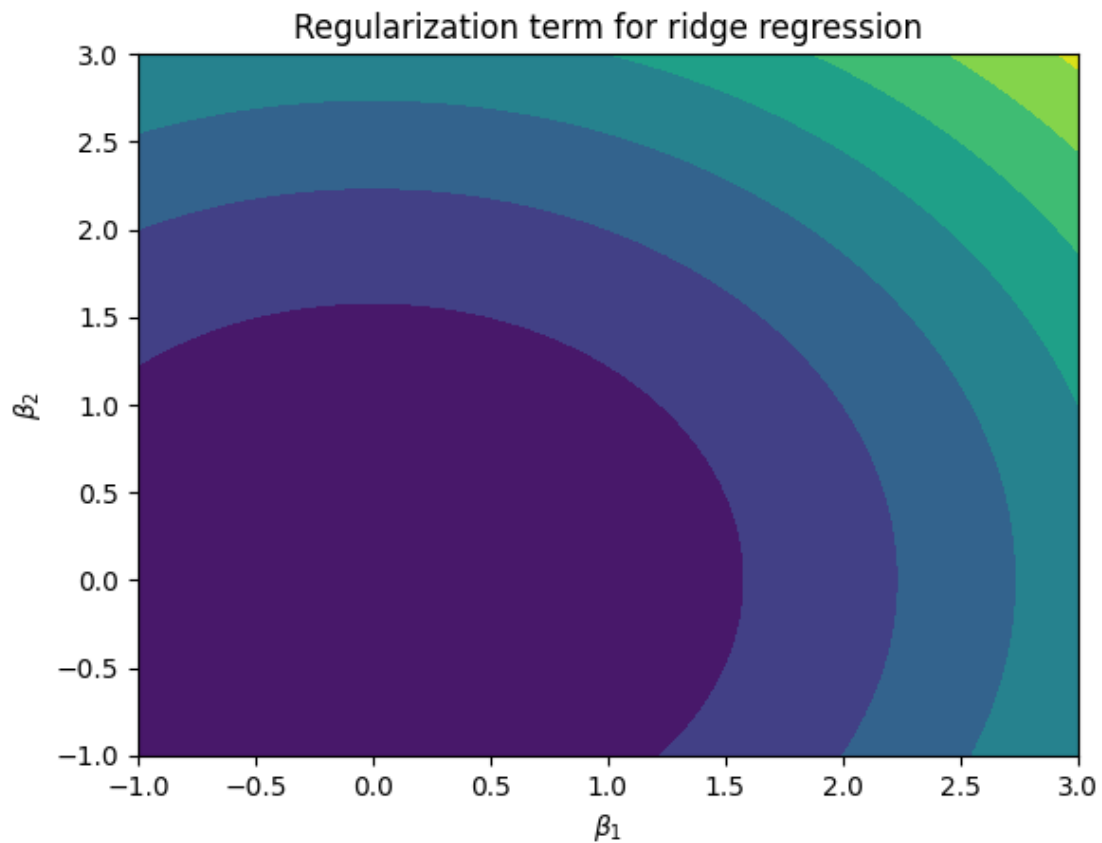
```

return np.abs(beta1)+np.abs(beta2)

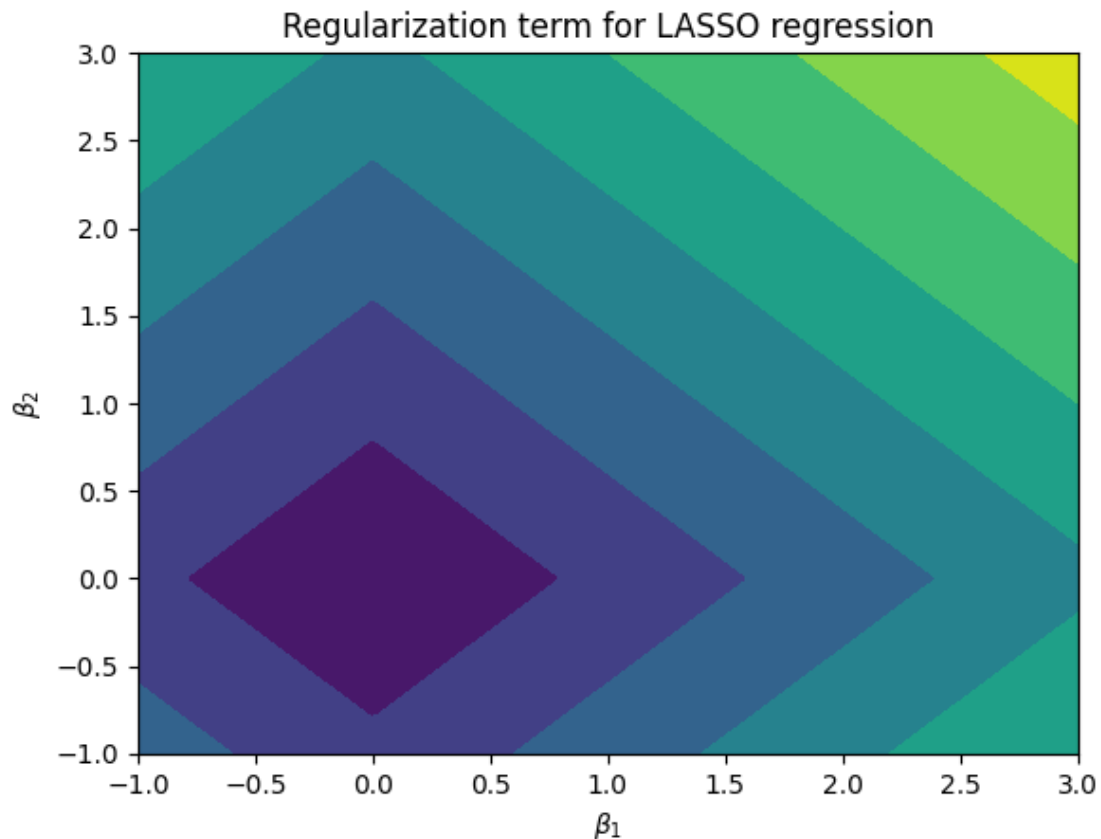
plt.title('Regularization term for ridge regression')
plt.xlabel(r'$\beta_1$')
plt.ylabel(r'$\beta_2$')
plt.contourf(beta1, beta2, regulariz_ridge(beta1,beta2))
plt.show()

plt.title('Regularization term for LASSO regression')
plt.xlabel(r'$\beta_1$')
plt.ylabel(r'$\beta_2$')
plt.contourf(beta1, beta2, regulariz_lasso(beta1,beta2))

```



[7]: <matplotlib.contour.QuadContourSet at 0x7f2116bd3a00>

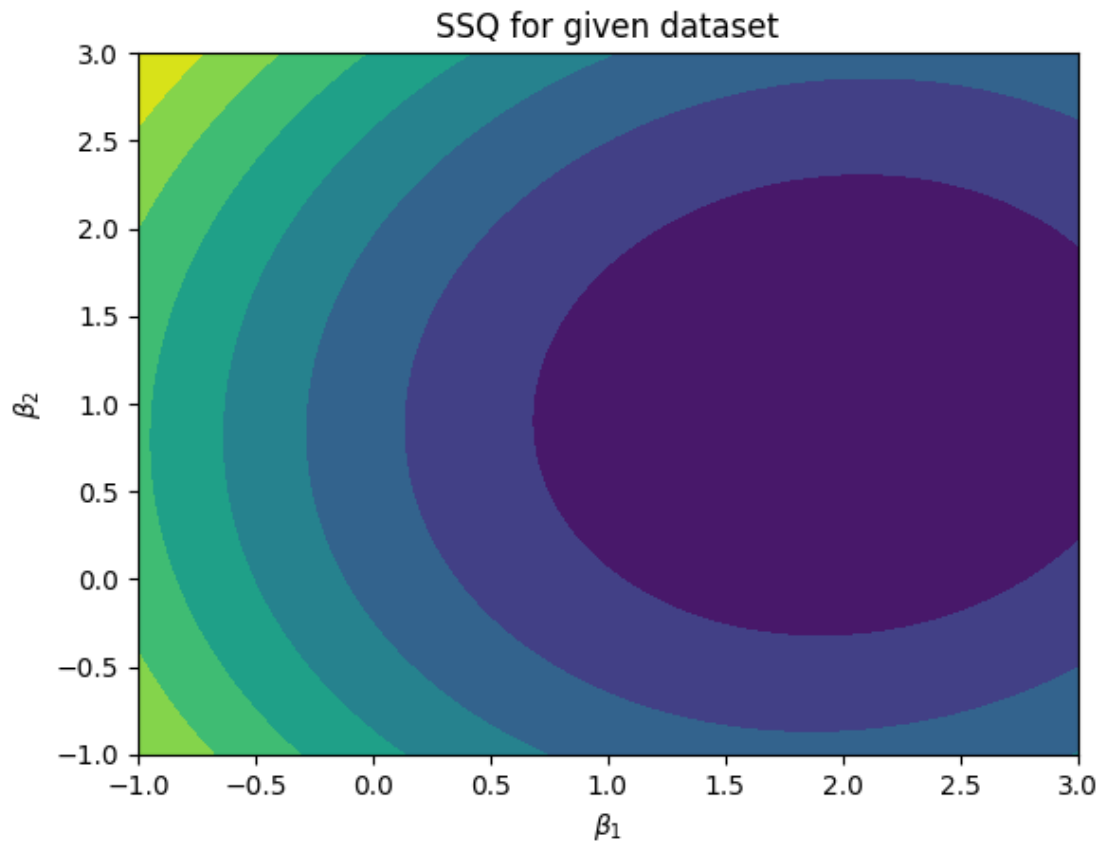


```
[8]: # TODO: for each combination of parameters, compute the sum of squared
      ↪ deviations.
      # do not use loops, but numpy broadcasting!
      # TODO: make a contour plot for sum of squared deviations

      # the stuff below uses loops, can't figure out how to do it with broadcasting :/
      def SSQ(beta1,beta2,features, labels):
          z = np.zeros((beta1.shape[0], beta2.shape[0]))
          for i,b1 in enumerate(beta1):
              for j,b2 in enumerate(beta2):
                  beta = [b1, b2]
                  z[i,j] = np.sum((beta @ features - labels)**2)
          return z

      #print(np.outer(beta1,x[0,:]).shape, beta1.shape, x[0,:].shape)
      #print(beta1)
      #print(beta1.shape, beta2.shape, z.shape)
      beta1 = np.linspace(-1,3,150)
      beta2 = np.linspace(-1,3,150)
      z = SSQ(beta1, beta2, x, y)
```

```
plt.contourf(beta1,beta2,z)
plt.title('SSQ for given dataset')
plt.xlabel(r'$\beta_1$')
plt.ylabel(r'$\beta_2$')
plt.show()
```



(c)

```
[9]: # TODO: for each lambda, plot both ridge regression and lasso loss functions
def loss(beta1, beta2, lam, features, labels, method='ridge'):
    mbeta1, mbeta2 = np.meshgrid(beta1, beta2) # use meshgrid-functions for
    ↪ regulariz.terms
    regulariz = np.zeros((beta1.shape[0], beta2.shape[0]))
    if method=='ridge':
        regulariz = regulariz_ridge(mbeta1, mbeta2)
    elif method=='lasso':
        regulariz = regulariz_lasso(mbeta1, mbeta2)
    else:
```





```


        raise NameError(f'Name {method} is not a specified regularization_
↪method :(')
    return SSQ(beta1, beta2, features, labels) + lam*regulariz


for meth in ['ridge', 'lasso']: # lets _speed_ this up with some _meth_s
    for lam in [0, 10, 50, 100, 200, 300]:
        lss = loss(beta1, beta2, lam, x, y, method=meth)
        plt.contourf(beta1, beta2, lss)
        plt.colorbar()
        plt.title(f'{meth} loss for given dataset, $\lambda$={lam}$')
        plt.xlabel(r'$\beta_1$')
        plt.ylabel(r'$\beta_2$')
        plt.show()

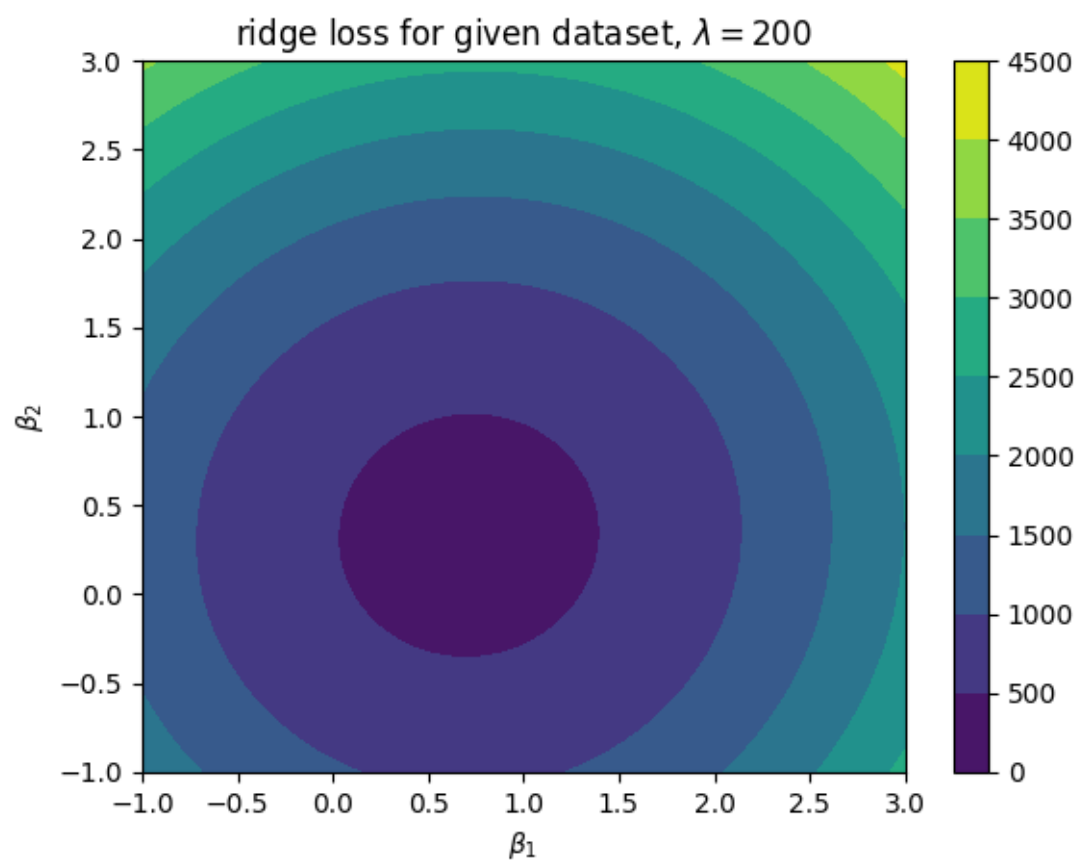
```










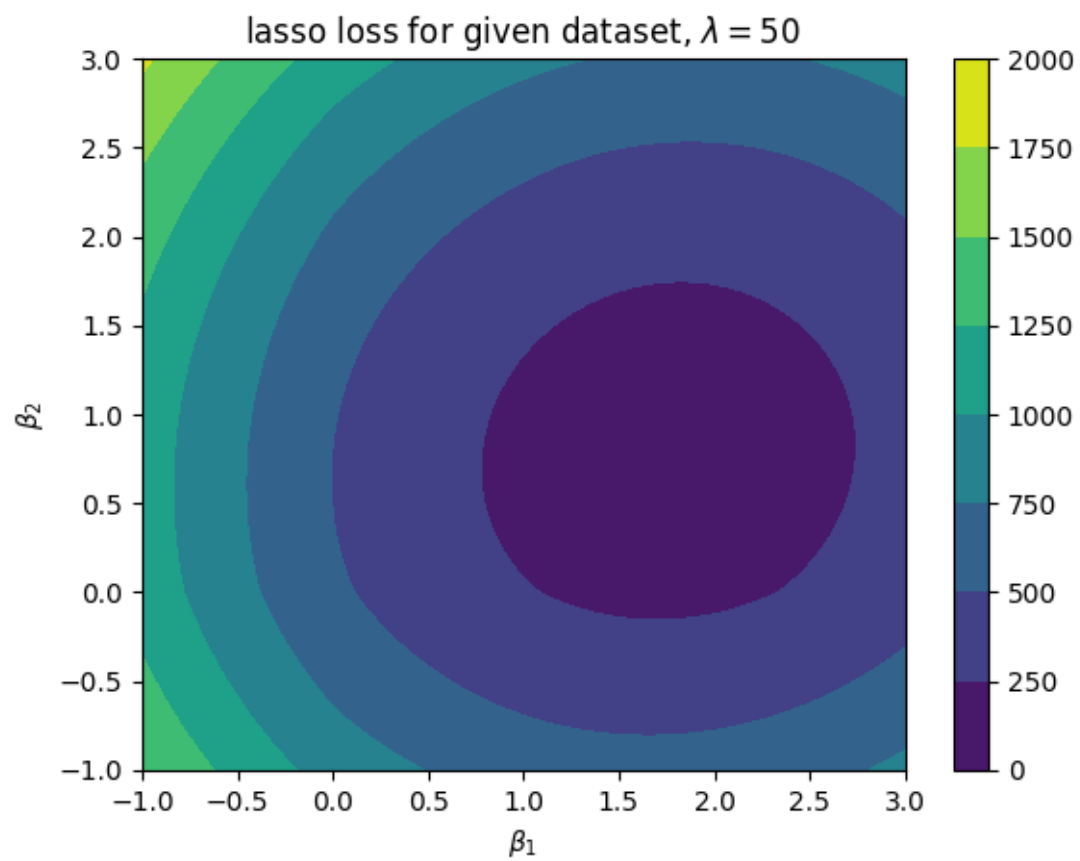





















Discussion: We observe the following:

- For higher values of λ , we see that the overall magnitude of losses increases (obviously), hence the gradient of the loss surfaces is increased as well.
- For higher values of λ , the shape of the loss surface stays roughly the same for ridge regression: It is transformed from a elliptical shape to a more circular shape, i.e. the eccentricity of the contour lines shrinks.
- For ridge regression, higher values of λ distort the contour lines towards a rhombus: The loss surfaces are just a superposition of the (elliptical) SSQ and the (rhomboidal) lasso loss.

1.5 5 CT reconstruction

set up design matrix (run this once to save to disk)

```
[10]: # create design matrix
# don't change any of this, just run it once to create and save the design
      ↪matrix
import os

if not os.path.exists('data/design_matrix.npy'):
```

```

res = (99, 117)
xs = np.arange(0, res[1]+1) - res[1]/2 # np.linspace(-1, 1, res[1] + 1)
ys = np.arange(0, res[0]+1) - res[0]/2 #np.linspace(-1, 1, res[0] + 1)

# rays are defined by origin and direction
n_parallel_rays = 70
ray_offset_range = [-res[1]/1.5, res[1]/1.5]
n_ray_angles = 30
n_rays = n_parallel_rays * n_ray_angles

ray_angles = np.linspace(0, np.pi, n_ray_angles, endpoint=False) + np.pi/
↪n_ray_angles

# offsets for ray_angle = 0, i.e. parallel to x-axis
ray_0_offsets = np.stack([np.zeros(n_parallel_rays), np.
↪linspace(*ray_offset_range, n_parallel_rays)], axis=-1)
ray_0_directions = np.stack([np.ones(n_parallel_rays), np.
↪zeros(n_parallel_rays)], axis=-1)

def rot_mat(angle):
    c, s = np.cos(angle), np.sin(angle)
    return np.stack([np.stack([c, s], axis=-1), np.stack([-s, c],
↪axis=-1)], axis=-1)

ray_rot_mats = rot_mat(ray_angles)

ray_offsets = np.einsum('oi,aij->aoj', ray_0_offsets, ray_rot_mats).
↪reshape(-1, 2)
ray_directions = np.einsum('oi,aij->aoj', ray_0_directions, ray_rot_mats).
↪reshape(-1, 2)

sigma = 1
kernel = lambda x: np.exp(-x**2/sigma**2/2)

xsc = (xs[1:] + xs[:-1]) / 2
ysc = (ys[1:] + ys[:-1]) / 2
b = np.stack(np.meshgrid(xsc, ysc), axis=-1).reshape(-1, 2)
a = ray_offsets
v = ray_directions
v = v / np.linalg.norm(v, axis=-1, keepdims=True)
p = ((b[None] - a[:, None]) * v[:, None]).sum(-1, keepdims=True) * v[:,
↪None] + a[:, None]
d = np.linalg.norm(b - p, axis=-1)
d = kernel(d)
design_matrix = d.T

```

```

np.save('data/design_matrix.npy', design_matrix)
print(f'created and saved design matrix of shape {design_matrix.shape} at_
↳data/design_matrix.npy')

```

(a)


```

[11]: design_matrix = np.load('data/design_matrix.npy')
res = (99, 117)

# TODO: visualize four random columns as images, using an image shape of (99,
↳117)
img_shape = (99, 117)

fig, axs = plt.subplots(1, 4, figsize=(16, 4))
for i, ax in zip(np.random.choice(np.arange(design_matrix.shape[1]), 4), axs):
    ax.imshow(design_matrix[:, i].reshape(*res));
    ax.set_title(f'column {i}')

```




Interpretation of a column of X : Every column of X describes the axis along which the detectors are positioned relative to the object to be reconstructed. Weighing the original object with a column of X gives the part of the object which can be projected onto the detector, the integral over this then gives the detector signal strength.

```

[12]: sino = np.load('data/sino.npy')

# visualize sinogram as image
n_parallel_rays = 70
n_angles = 30
plt.imshow(sino.reshape(n_angles, n_parallel_rays), origin='lower')
plt.show();

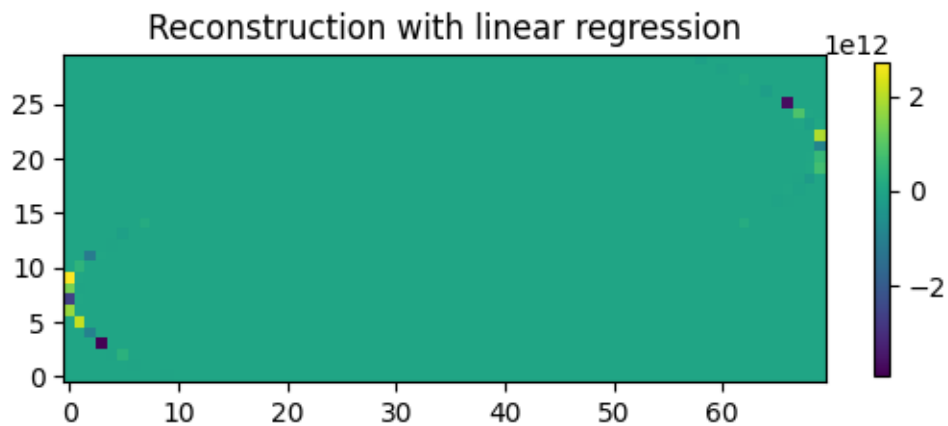
```



(b)

```
[13]: # TODO: solve the reconstruction with linear regression and visualize the result
X = design_matrix
beta = np.linalg.inv(X @ X.T) @ X @ sino.T
```

```
[14]: reconst = beta.T @ X
plt.imshow(reconst.reshape(n_angles, n_parallel_rays), origin='lower')
plt.colorbar(shrink=.44)
plt.title('Reconstruction with linear regression')
plt.show()
```




```
[15]: # TODO: solve the reconstruction with ridge regression and visualize the result
# Optional: try out different regularization strengths and observe the influence
X = design_matrix
```


```
#beta_r = np.linalg.inv(X @ X.T + lam*np.identity(X.shape[0])) @ X @ sino.T
```

```
[16]: for lam in [10, 50, 100, 200, 300]:
    beta_r = np.linalg.inv(X @ X.T + lam*np.identity(X.shape[0])) @ X @ sino.T
    reconst = beta_r.T @ X
    plt.subplot(1,2,1)
    plt.title('Reconstructed image')
    plt.imshow(reconst.reshape(n_angles, n_parallel_rays), origin='lower')
    plt.colorbar(shrink=0.22)
    #plt.show()
    plt.subplot(1,2,2)
    plt.title('Difference to original')
    plt.imshow((reconst-sino).reshape(n_angles, n_parallel_rays),
    ↪origin='lower')
    plt.colorbar(shrink=0.22)
    plt.suptitle(f"Results of Ridge regression with  $\lambda={lam}$ ", y=0.7,
    ↪fontsize=14)
    plt.show()
```


Results of Ridge regression with $\lambda = 10$




Results of Ridge regression with $\lambda = 50$




Results of Ridge regression with $\lambda = 100$



Results of Ridge regression with $\lambda = 200$



Results of Ridge regression with $\lambda = 300$



We observe that:

- The plain Linear regression fails phenomenally at the task of reconstruction the image, because of the very inaccurate matrix inversion
- For all values of $\lambda \in \{10, 50, 100, 200, 300\}$, the ridge regression give optically fine results, for the smallest value of $\lambda = 10$, we get the least deviation from the original image.
- The shape of the difference does not vary by a lot, only its magnitude changes meaningfully with λ